

TEAM BRIEFING · INTERNAL

CareerConnect Job Portal

Launch Plan

Plan of action, known issues, owners and deadlines for the public launch on **Sunday 11 October 2026**.

15 days

26 Sep → 11 Oct 2026

6 people

Ram, Vivek, Imran, Sneha,
Yug, Tripti

54 tasks

33 are P0 (must ship)

₹0

Free tiers only

Prepared by	Ram Mohan Dixit, Manager
Date	26 September 2026
Code baseline	<code>origin/main</code> at <code>2c84c10</code> (23 Sep 2026). All 9 remote branches reviewed.
Working branch	<code>launch/job-portal</code>
Live tracker	<code>docs/launch/Launch_Tracker.csv</code> (the source of truth; this PDF is a snapshot)

How to read this: Sections 1–3 explain why we are doing this and what is wrong today. Sections 4–6 cover the plan and each person's tasks. Sections 7–10 cover job supply, infrastructure, the launch gate and the revenue plan. The appendix has the credential-rotation runbook and team working rules.

1. Summary

We launch **only the job portal** on **11 October 2026**. Courses, payments, assessments and AI interviews stay in the code but are switched off. We spend **no money**, so everything runs on free tiers.

Three things decide whether the launch succeeds:

1. **Security and trust.** Today an employer can publish a job without admin approval, anyone can register as an employer and post straight away, live updates are broadcast to every connected browser, and a logged-in user can unlock a paid course without paying. Scam jobs or leaked candidate data would end us on day one.
2. **Real job supply.** The live site shows 95 jobs, and **91 of them are scraped**, mostly from LinkedIn and Internshala, which likely breaks their terms, and some with invented salaries. We need **200+ real, approved listings** by 10 October.
3. **Stability on free hosting.** Right now the backend can serve about 20–40 people at the same moment. After the free fixes in this plan it should reach about 100 concurrent and 1,500–3,000 daily users.

Decision: the architecture is right for our stage (one React frontend, one Express API, MongoDB), so **no rewrite**. The work is enforcing rules on the server, moving heavy work out of requests, and closing launch gaps.

Launch scope

In scope (launching)

- Job and internship discovery, filters and search
- Public job and internship detail pages (new)
- Candidate signup (email OTP, Google), profile, resume builder and upload
- Apply, track and withdraw applications
- Employer registration (with approval), job posting, applicant pipeline
- Human interviews, scorecards, offers
- Notifications and emails (new)
- Admin moderation of employers and jobs

Out of scope (hidden with feature flags)

- Courses, the learning hub, employee training
- Payments (Razorpay)
- Assessments (builder, hub, candidate room)
- AI interviews
- Messaging (Imran's branch)
- LinkedIn and Internshala scraping

These are postponed, not dropped. Several of them become paid features later (section 10).

Key dates

Date	Milestone
26–27 Sep	Step 0: branches set up (done), hosting confirmed (done), credentials rotated, kickoff call
30 Sep	Phase 1 gate: no known way to self-publish, get a free paid item, or read another user's data
5 Oct	Phase 2 gate: all launch blockers closed
6 Oct, evening	Code freeze. <code>launch/job-portal</code> is merged to <code>main</code> . Only bug fixes after this.
6–8 Oct	Closed beta: 50 students and 10 employers. Full regression and load test.
10 Oct, evening	Go/no-go meeting
11 Oct	Public launch
11–25 Oct	Close support after launch: triage at 10:00, metrics review at 18:00

2. Where we are today

2.1 How many users can we handle?

This assumes free tiers: Render Free (confirmed) (0.1 CPU, 512 MB RAM, sleeps after 15 minutes idle), MongoDB Atlas M0, Cloudinary Free, Brevo Free and Gemini Free. **These are estimates from reading the code and timing a few live requests. We haven't load-tested yet;** Tripti's task Q03 confirms the numbers. One real measurement: the live job list took 0.7–4.3 seconds for one user.

Measure	Today	After the Phase 1–2 fixes (still ₹0)
People using the site at the same moment	about 20–40	about 100
Active users per day	about 300–500	about 1,500–3,000
Email signups per day	about 250–300 (Brevo free cap is 300 emails/day)	Same cap. Add the SMTP fallback; Google sign-in isn't affected.
First visitor after 15 minutes idle	30–60 second wait (cold start)	No wait (UptimeRobot keep-alive)
Link shared in a big WhatsApp group	Slow pages and timeouts	Holds up

What limits us most, in order:

1. Live scraping on every job or internship list request.
2. Password hashing: bcrypt cost 12, in pure JavaScript, on 0.1 CPU.
3. The 300 emails/day cap.
4. AI resume routes with no rate limit, plus a retry chain of up to 14 Gemini calls per click.
5. No compression and no caching.
6. A full scan of every interview each minute.

If we outgrow the free tiers: Oracle Cloud Always Free (4 ARM cores, 24 GB RAM, ₹0) for the API, and Cloudflare Pages for the frontend. Vercel Hobby doesn't allow commercial use, so we must move the frontend before we charge money.

2.2 System design scorecard

Area	Score	Why
Product and domain design	8/10	Rich, well-thought-out hiring flow: jobs → applications → interviews → offers
Architecture choice	7.5/10	A single API with one frontend and MongoDB is correct at our scale. Microservices would slow us down.
Security design	4/10	The server trusts values the client sends, and auth has four overlapping methods
Scalability	4.5/10	Heavy work runs inside requests, the in-memory cache and sockets won't work across servers, and there are timers in the API process
Reliability and operations	4/10	No backups on M0, no monitoring, many errors silently swallowed
Maintainability	5/10	Controllers of 2,000–3,600 lines, and about 9,000 lines of profile code copied three times
Overall	about 5.5/10	A strong MVP that needs hardening. It can be fixed without a rewrite.

3. Known issues (checked on latest `main`)

Each issue shows the tracker task that fixes it. Severity: **Critical** means we can't launch with it, **High** means fix before launch, and **Medium** means fix if time allows.

Severity	Issue	Where	Task
Critical	A paid course can be unlocked without paying. Any order ID starting with <code>order_test_</code> , or a request with no signature, passes the payment check. Razorpay test key and secret are also hardcoded as fallbacks.	<code>server/controllers/paymentController.js</code>	S02, Z03
Critical	Live updates leak data to everyone. Sockets need no login, allow any origin, let anyone join any room, and <code>io.emit</code> sends full interview and application records to every connected browser.	<code>server/services/socketService.js</code>	S05
Critical	Admin approval never happens. <code>createJob</code> takes <code>status</code> from the request, the employer form sends "Published", and the editable-field allowlist still includes <code>status</code> . Internships behave the same.	<code>jobController.js</code> , <code>internshipController.js</code> , <code>utils/listingSecurity.js</code> , <code>JobModal.jsx</code>	S03
Critical	Anyone can register as an employer and post straight away. There's no verification or approval, and free-mail domains are allowed.	<code>authController.registerEmployer</code> , <code>roleMiddleware.js</code>	S04
Critical	Leaked credentials. <code>server/.env</code> (MongoDB connection string and JWT secret) is still in git history from 23 Aug.	git commits <code>a110544</code> , <code>3214bdc</code>	Z03
High	No Privacy Policy or Terms (the footer links go to <code>/home</code>), no consent at signup, and no account deletion. We store resumes, which are personal data under the DPDP Act 2023.	<code>client/src/pages/Home.jsx</code>	G05, G07
High	Scraped jobs: LinkedIn and Internshala HTML scraping likely breaks their terms, salaries and responsibilities are invented, IDs change on every request, and scraping runs per request.	<code>services/jobScraperService.js</code> , <code>jobController.js</code>	S07, I04
High	Rate limiting is mostly off: <code>globalLimiter</code> , <code>otpLimiter</code> and <code>aiResumeLimiter</code> do nothing. <code>trust proxy</code> isn't set, so behind Render every user shares one IP-based limit.	<code>middleware/rateLimitMiddleware.js</code> , <code>app.js</code>	S06
High	No job detail page with its own URL (can't share or index a job). The employer link <code>/jobs/:id/applications</code> is a 404.	<code>client/src/App.jsx</code>	G01
High	Candidates get no email when their application status changes, an interview is scheduled, or an offer is made.	<code>applicationController.js</code> , <code>offerController.js</code>	G03

High	Admin settings (auto-approve, allow employer registration, maintenance mode) are saved but never checked.	models/PlatformSetting.js	S04
High	Auth has four overlapping methods: cookie JWT, a Bearer token in localStorage (an XSS risk), a separate admin JWT without the av version claim, and unused session code.	authMiddleware.js, adminController.js, client/src/api/api.jsx	S09
High	Missing role checks: student routes only check login, employers can apply to jobs, and a duplicate apply path trusts the email the client sends.	studentRoutes.js, applicationController.js	S08
High	Jobs never expire after their deadline, and no React error boundary exists (one crash blanks the app).	server.js, client/src/App.jsx	G04, G06
Medium	AI resume routes have no limit, and the retry chain makes up to 14 Gemini calls per click. Every resume edit uploads a new Cloudinary file and delete never removes it.	aiResumeservice.js, resumeController.js	G09, G08
Medium	Performance: no compression, no cache headers, search runs a regex full scan, no {status, createdAt} index, and the per-minute interview sweep loads every document.	app.js, models/Job.js, server.js	I06, I07
Medium	No backups on Atlas M0, no monitoring or error tracking, and many empty catch {} blocks.	infrastructure	I02, I08
Medium	Maintainability: adminController is 3,644 lines, resumeController 3,096 and EmployerDashboard.jsx 2,251. About 9,000 lines of profile code are copied across student, fresher and professional. After launch.	controllers, profile components	Post-launch

Parked work with known bugs (feature/assessments-ai-interview, not launching): the interview sweep marks unstated AI interviews "completed", the assessment timer can be bypassed by submitting without an attemptId, answers are lost on refresh, the "Assign Assessment" button does nothing, and dictation can write into the wrong answer. Fix all of these before that branch is merged.

4. Branch status and rules

Checked on 26 Sep against `origin/main` at `2c84c10`. Each merge was tried in a throwaway copy of the repo, so no branch was changed.

Branch	Work not in main	Merges cleanly?	Decision
<code>launch/job-portal</code>	Launch plan docs	n/a	All launch work goes here. Tripti reviews every PR.
<code>ram-mohan-code</code>	6 commits: ATS scoring, JD upload, verified ATS PDF workflow	No: <code>.gitignore</code> , <code>ATSCheckerAndFixer.jsx</code> , <code>ResumeBuilder.jsx</code> , <code>vite.config.js</code>	Ram and Yug decide by 29 Sep whether it ships (B01)
<code>Imran</code>	1 commit (about 5,300 lines): messaging, job visibility, offer lifecycle, CIN/GST field	No: 4 files	On hold, out of scope. New chat rooms need auth first (B02).
<code>sneha</code>	Course page and course controller, plus one auth fix	No: 3 files	Cherry-pick only the auth fix (B03)
<code>feature/assessments-ai-interview</code>	Parked assessments and AI interview work	No: 12 files	Don't merge before launch
<code>develop</code> , <code>tripti</code> , <code>vivek</code> , <code>yug</code>	None (already in main)	n/a	Stale. Don't open PRs from them (B04).

Working rules until launch

- Branch from `launch/job-portal`. One small PR per tracker task, with the ID in the title, for example `[S03] Force pending approval on new jobs`.
- Every PR needs one reviewer (the Reviewer column in the tracker), a green CI run and a short "how I tested" note.
- Security fixes (S02–S05, S08) must include an automated test that proves the hole is closed.
- **Nobody merges into main before code freeze.** Render deploys `main` to production automatically.
- **No new features after code freeze on the evening of 6 Oct.** Only fixes for bugs found in QA or the beta.
- Never commit `.env` files or keys. Secrets live only in the Render and Vercel dashboards.
- Update your rows in the tracker's `Status` column every evening: Not started, In progress, In review or Done.

5. The 15-day plan

Phase	Dates	Goal	Exit gate
Step 0: Set up	26–27 Sep	Branches set up (done), credentials rotated, kickoff call, free-tier plans confirmed	New secrets live; everyone working from launch/job-portal
P1: Lock down and cut scope	27–30 Sep	Payment bypass, socket lockdown, moderation, employer verification, feature flags, rate limits, scraping off, backups, keep-alive	No known way to self-publish, get a free paid item, or read another user's data
P2: Close the gaps	1–5 Oct	Public job pages and SEO, candidate emails, expiry, legal pages and consent, account deletion, scheduled job feeds, performance, Sentry	Every P0 task done
P3: Freeze and test	6–8 Oct	Code freeze, merge to main, full regression, k6 load test on staging, closed beta	Load test meets target; no open P0 bugs
P4: Fix and decide	9–10 Oct	Fix beta bugs; go/no-go meeting	Checklist in section 9 is all ticked
Launch and close support	11–25 Oct	Launch, then triage at 10:00 and metrics at 18:00 every day	—

Daily rhythm (everyone)

- **10:00:** 15-minute standup. What you did yesterday, what you'll do today, and what's blocking you.
- **1 hour a day:** employer outreach, 10 personal messages each (section 7). Job supply is our biggest risk, so everyone helps.
- **Evening:** update the tracker Status column. Ram posts a short progress note in the team group.

Who owns what

Person	Area	Main deliverables
Ram	Manager, product, supply	Scope, moderation rules, legal copy, 50+ employers and 200+ real jobs, Geeta University pilot, launch channels, support inbox, pricing interviews
Vivek	Backend	Payment bypass (S02), moderation (S03), employer verification (S04), socket lockdown (S05), role checks (S08), admin token (S09), bcrypt (S10), emails (G03), expiry (G04), interview sweep (I07)
Imran	Infrastructure	Credential rotation (Z03), rate limits and proxy (S06), scraping off (S07), keep-alive (I01), backups (I02), Cloudflare (I03), scheduled job feeds (I04, I05), performance (I06), Sentry (I08)
Sneha	Candidate frontend	Feature flags (S01), public job pages (G01), SEO (G02), legal pages and consent (G05), error boundary (G06), report-job button (G11), public-path fix (G12)
Yug	Resume	AI limits (G09), Cloudinary cleanup (G08), account deletion (G07), resume privacy and DOCX handling (G10)
Tripti	QA	Test checklist (Q01), regression tests (Q02), load test (Q03), code freeze (Q04 with Vivek), beta (Q05), go/no-go (Q06)

6. Task list by owner

Snapshot of `docs/launch/Launch_Tracker.csv` on 26 Sep: 54 tasks, 3 done. The CSV is the live version. Priority: **P0** must ship, **P1** should ship, **P2** nice to have.

Ram · Manager: product, scope, legal copy, employer supply, launch marketing · 12 tasks (7 P0)

ID	Pri	Task	Due	Depends	Done when
Z01	P0	Park uncommitted assessments/AI interview work on feature branch	26 Sep		feature/assessments-ai-interview pushed with known-issues list
Z02	P0	Create launch/job-portal branch from latest main	26 Sep		Branch pushed; team told to branch from it
Z04	P0	Team kickoff call (30 min): walk through README and tracker; confirm owners	27 Sep		Every owner has acknowledged their rows
B01	P0	Decide + merge ATS resume work (origin/ram-mohan-code, 6 commits: ATS scoring, JD upload, verified ATS PDF workflow) into launch/job-portal; conflicts in .gitignore, ATSCheckerAndFixer.jsx, ResumeBuilder.jsx, vite.config.js	29 Sep		PR merged into launch/job-portal, unit tests pass, client builds; or explicitly deferred
L01	P0	Moderation SOP: what verifies an employer, banned job types (any fee charged to candidates), response time, admin rota	30 Sep		Written SOP shared; rota for 11-25 Oct assigned
J02	P0	Geeta University placement cell: recruiter list + alumni companies; university as pilot partner	05 Oct		Meeting done; recruiter list received; pilot agreed
L03	P1	Support inbox + FAQ + canned replies	08 Oct		Support email monitored daily
L02	P1	Launch channels: WhatsApp/Telegram jobs channel, LinkedIn page, college ambassadors, launch posts drafted	09 Oct		Channels live; launch posts scheduled
J03	P1	Assisted posting: team posts jobs from JDs employers send on WhatsApp/email (with employer consent)	10 Oct	S04	Each assisted post has employer contact and consent recorded
J04	P1	Government jobs daily curation process (Employment News / NCS / state PSC notifications)	10 Oct	G01	Daily posts with official source link
L04	P2	Pricing discovery: interview 10 employers on featured jobs / paid plans	10 Oct		Notes from 10 interviews; draft pricing
Q06	P0	Go/no-go meeting against README checklist	10 Oct	Q05	All checklist items ticked or explicitly accepted

Vivek · Backend: moderation, auth, emails, sockets, expiry · 11 tasks (7 P0)

ID	Pri	Task	Due	Depends	Done when
S02	P0	Remove Razorpay hardcoded fallbacks and order_test_/no-signature bypass in paymentController	28 Sep	Z03	No hardcoded keys; verify rejects unsigned/sandbox orders; test added
S03	P0	Enforce moderation: new jobs/internships always Pending Approval; employers cannot set Published (create + update + status endpoints + listingSecurity allowlist)	29 Sep		Integration test proves employer cannot self-publish; admin approve publishes
S10	P1	Lower bcrypt cost 12 -> 10 (existing hashes still verify)	29 Sep		Login CPU time roughly quartered; old accounts still log in
S04	P0	Employer verification: new employers blocked from posting until admin approves; wire PlatformSetting autoApproveJobs/autoApproveEmployers/allowEmployerRegistration	30 Sep	S03	Unapproved employer gets 403 on post; admin settings change behaviour
S05	P0	Lock down socket.io: authenticate handshake via JWT cookie, only join own rooms, remove io.emit broadcasts, restrict CORS origins	30 Sep		Second browser cannot receive another user's interview/application events

S08	P1	Role checks: student routes require student type; only candidates can apply; remove/secure /api/student/apply duplicate path	30 Sep	Employer apply returns 403; tests added
S09	P1	Fix admin JWT missing av claim and admin cookie sameSite for cross-site deploy	30 Sep	Admin stays logged in on production domain after password reset of another admin
G03	P0	Candidate emails for application status change, interview scheduled/changed, offer sent (respect 300/day Brevo; SMTP fallback)	03 Oct	Each event sends one email; daily volume counter logged
G04	P0	Auto-close jobs/internships after deadline (scheduled job) and block applying to expired listings	03 Oct	Expired listing disappears from lists and apply returns 400
I07	P1	Move interview sweep out of per-minute full scan (filter by time + index + lean) or to a scheduled worker	03 Oct	Sweep query uses index; no full collection load
Q04	P0	Code freeze: merge launch/job-portal into main; only bug fixes after	06 Oct	Freeze announced; main deployed to production

Imran · Infrastructure: secrets, Render, Cloudflare, backups, monitoring, job feeds, DB indexes · 12 tasks (7 P0)

ID	Pri	Task	Due	Depends	Done when
I01	P0	UptimeRobot (free) ping /health every 5 min; alert to team email. Required: Render Free sleeps after 15 min idle	27 Sep	Z05	No cold starts; downtime alert received in test
Z03	P0	Rotate MongoDB / JWT / SESSION / OTP secrets and Razorpay test key (see README runbook)	27 Sep		New values live on Render; old values confirmed dead; logged in security-hardening.md
Z05	P0	Confirm Render plan and deploy branch (confirmed 26 Sep: Render Free; production deploys from main)	27 Sep		Tiers written in README
B02	P1	Imran branch (1 commit, 5.3k lines: messaging, job visibility/eligibility, offer lifecycle, CIN/GST field): out of launch scope; hold. Before any merge, new conversation_ socket rooms need auth (same leak as S05) and hardcoded Geeta University fallback removed	28 Sep	S05	Decision recorded; if merged, only after S05 and with conflicts (App.jsx, CandidateInterviewsView.jsx, vite.config.js, offerController.js) resolved
S07	P0	Turn off LinkedIn and Internshala scraping; remove invented salary/responsibilities for external jobs	28 Sep		No LinkedIn/Internshala requests in logs; external jobs show only source data
S06	P0	Enable globalLimiter/otpLimiter/aiResumeLimiter; set TRUST_PROXY_HOPS on Render; rate-limit /api/internships/live	29 Sep		Limits verified per-IP from two different IPs
I02	P0	Nightly mongodump backup via GitHub Actions to private storage; test a restore	30 Sep	Z03	Restore into scratch DB succeeds
I03	P1	Put Cloudflare free proxy in front of the domain (DNS + SSL + caching of static assets)	30 Sep		Domain served via Cloudflare; site works end to end
I06	P1	compression middleware, Cache-Control on public lists, {status,createdAt} index, text index for search, lean() on application lists	03 Oct		Job list response gzipped; explain() shows index use
I04	P0	Replace per-request scraping with scheduled sync into DB (stable IDs, isExternal, source + attribution, expiry); admin-only sync trigger	04 Oct	S07	Job list served from DB only; p95 under 1.5 s
I05	P1	Add legal job feeds: Greenhouse / Lever / Ashby public boards for a curated list of Indian companies; Adzuna/Jobble/Careerjet (India) after checking terms	05 Oct	I04	500+ real external jobs with working apply links
I08	P1	Sentry free (client + server) and a free analytics tool	05 Oct		Test error visible in Sentry; page views tracked

Sneha · Candidate frontend: feature flags, job pages, SEO, legal pages, error boundary · 8 tasks (4 P0)

ID	Pri	Task	Due	Depends	Done when
----	-----	------	-----	---------	-----------

B03	P1	Sneha branch (course page + course controller changes): out of launch scope; cherry-pick only the authMiddleware fix that treats literal 'null'/'undefined' Bearer tokens as missing	28 Sep		Auth fix on launch/job-portal with a test; course changes parked
G12	P1	Fix public-path 401 redirect list (jobs/internships/organizations) in client api.jsx	29 Sep		Logged-out visitor never bounced to /login on public pages
S01	P0	Feature flag: hide courses, payments, assessments, AI interview (client VITE_ENABLE_* + server conditional route mounts)	29 Sep		No course/payment/assessment UI anywhere for any role; /api/courses, /api/course-content, /api/payment, /api/employer/learning, /api/assessments return 404
G06	P0	Global React error boundary + lazy-chunk load retry	01 Oct		Thrown error shows fallback UI instead of blank page
G01	P0	Public job and internship detail pages at /jobs/:id and /internships/:id (no login to view), fix employer /jobs/:id/applications 404	03 Oct	S03	Shareable link opens for logged-out visitor; apply prompts login
G05	P0	Privacy Policy, Terms, Contact pages; fix footer links; consent checkbox on all signup paths (DPDP notice + grievance contact)	04 Oct	L01	Pages live; signup blocked without consent; consent timestamp stored
G02	P1	SEO: per-page title/description/OG tags, JobPosting JSON-LD on employer-posted jobs, sitemap.xml, robots.txt	05 Oct	G01	Google Rich Results Test passes for a job page; sitemap lists live jobs
G11	P1	Report this job button + endpoint feeding admin reports	05 Oct		Report appears in admin reports within one minute

Yug · Resume: AI limits, Cloudinary cleanup, account deletion, resume privacy · 4 tasks (2 P0)

ID	Pri	Task	Due	Depends	Done when
G08	P1	Cloudinary cleanup: destroy old asset on resume update/delete	30 Sep		Asset count stable after 5 edits of one resume
G09	P0	Rate-limit AI resume routes; cap Gemini retry chain to 2 attempts; per-user daily free AI quota	30 Sep		Quota exceeded returns friendly 429; no more than 2 Gemini calls per action
G10	P1	Resume uploads: either parse DOCX or accept PDF only; verify only employers the candidate applied to can view a resume	03 Oct		Upload rules consistent; unauthorized resume URL returns 403
G07	P0	Account deletion (candidate + employer) that also deletes resumes from Cloudinary	04 Oct	G08	User can delete account; Cloudinary assets gone

Tripti · QA: test checklist, regression tests, load test, beta, go/no-go · 4 tasks (4 P0)

ID	Pri	Task	Due	Depends	Done when
Q01	P0	Four-role test checklist (visitor/candidate/employer/admin) covering launch scope only	30 Sep		Checklist reviewed by each owner
Q02	P0	Automated regression tests for S02/S03/S04/S05/S08 fixes	03 Oct	S02, S03, S04, S05, S08	Tests in CI and passing
Q03	P0	k6 load test against staging copy (not production DB)	07 Oct	I04, I06	100 concurrent users, job list p95 < 1.5 s, errors < 1%
Q05	P0	Closed beta: 50 students + 10 employers; bug log and daily triage	08 Oct	Q04	All beta P0 bugs fixed

All · Shared tasks · 3 tasks (2 P0)

ID	Pri	Task	Due	Depends	Done when
B04	P1	Personal branches are stale (yug 128, Imran/sneha 58, ram-mohan-code 57, tripti 12, vivek 4 commits behind main): everyone starts new work from launch/job-portal	27 Sep	Z04	No PR opened from a stale personal branch
J01	P0	Employer outreach: each member 10 personal messages/day (LinkedIn, WhatsApp, email); Founding Employer offer	10 Oct		50+ approved employers and 200+ real listings by 10 Oct

Q07	P0	Public launch + daily triage (10:00) and metrics review (18:00) until 25 Oct	25 Oct	Q06	Launched; daily reports posted
------------	-----------	--	--------	-----	--------------------------------

7. How we get jobs for ₹0

Once LinkedIn and Internshala scraping is off, we'll have very few Indian jobs. **Job supply is the number one launch risk**, so everyone spends 1 hour a day on it.

A. Real employers (most important)

- **Geeta University placement cell first:** get their recruiter list and alumni-run companies, and make the university our pilot partner (J02).
- **"Founding Employer" offer:** free job posts for life plus a badge for the first 100 companies.
- **Outreach target:** 6 people × 10 messages a day × 12 days is about 700 conversations, which should turn into roughly 50–80 employers (J01). Focus on startups, agencies and local businesses in Delhi NCR, Panipat and Haryana.
- **Assisted posting:** an employer WhatsApps us a JD and we post it for them, with their consent recorded (J03). Removing that friction matters more than any feature.

B. Public job feeds that are legal to use (Imran, I04/I05)

- **Careers pages on applicant-tracking systems** publish job boards meant for sharing: Greenhouse (`boards-api.greenhouse.io/v1/boards/<company>/jobs`), Lever (`api.lever.co/v0/postings/<company>`) and Ashby. A curated list of 200–300 Indian startups gives us thousands of real jobs with genuine apply links.
- **Aggregator APIs with free keys that cover India:** Adzuna, Jooble and Careerjet. Some also pay per click.
- Keep **Remote and Arbeitnow** (already built) for remote jobs, with attribution.
- **Rules:** check each provider's terms first, store results in our database on a schedule (never scrape per request), label them "External job", always link to the original, and never invent a salary.

C. Government jobs (J04)

A daily hand-curated post of government job notifications (Employment News, the National Career Service portal, state commissions), with the official source link. These are public information, and candidates search for them daily.

D. Candidates attract employers

- **Google for Jobs:** public job pages plus JobPosting structured data plus a sitemap bring free search traffic (G01, G02). Only mark up jobs posted directly on our site.
- A daily WhatsApp/Telegram jobs channel, our LinkedIn page, and student ambassadors at colleges (L02).

8. Infrastructure on free tiers

Service	Free tier limit to know	What we do
Render Free (API, confirmed)	0.1 CPU, 512 MB RAM, sleeps after 15 minutes idle; deploys from <code>main</code>	UptimeRobot ping to <code>/health</code> every 5 minutes (I01); stop scraping per request; compression
MongoDB Atlas M0	512 MB storage, shared CPU, no automatic backups	Nightly <code>mongodump</code> via GitHub Actions and one tested restore (I02); proper indexes (I06)
Cloudinary	25 credits/month (storage + bandwidth + transformations)	Delete the old file on resume update and delete (G08)
Brevo	300 emails/day	Only important emails; SMTP fallback; log daily volume (G03)
Gemini	Free-tier requests per minute and per day	Per-user daily AI quota; at most 2 attempts per action (G09)
Vercel Hobby (frontend)	Non-commercial use only	OK for the free launch. Move to Cloudflare Pages (free, commercial allowed) before we charge money.
Cloudflare	Free DNS, SSL, CDN, DDoS protection	Put it in front of our domain (I03)
Sentry and analytics	Free developer plans	Error tracking for client and server (I08)

Confirmed on 26 Sep: the API runs on **Render Free**, and production deploys from `main`. So the live code is `main` at `2c84c10`, the keep-alive (I01) is required, and **any merge into `main` goes straight to production**.

Rollback plan

Keep the last known-good deploy ID for both Render and Vercel. If a release breaks production, redeploy the previous build from the dashboard (about 2 minutes) and then fix forward on `launch/job-portal`. Rehearse this once before 10 Oct.

9. Go/no-go checklist (evening of 10 Oct)

We launch only if every item is ticked, or explicitly accepted as a known risk by the whole team.

- ☐ All P0 tasks in the tracker are Done.
- ☐ Credentials rotated, and the old ones confirmed not to work.
- ☐ 200+ real, admin-approved jobs and internships are live.
- ☐ LinkedIn and Internshala scraping is off, and no invented salaries appear anywhere.
- ☐ Employers can't self-publish; unapproved employers can't post (automated tests pass).
- ☐ No socket event reaches a user it wasn't meant for (tested with two browsers).
- ☐ Load test on staging: 100 concurrent users, job list p95 under 1.5 s, error rate under 1%.
- ☐ A backup has been restored successfully into a test database at least once.
- ☐ Privacy Policy, Terms and Contact pages are live; signup asks for consent; account deletion works.
- ☐ Courses, payments, assessments and AI interviews are invisible for every role.
- ☐ Admin moderation rota assigned for 11–25 Oct, with a support inbox monitored daily.
- ☐ Rollback rehearsed.

What we measure after launch

Metric	Target, first 2 weeks
Real (non-scraped) active listings	200 at launch, growing about 20 a day
Approved employers	50 at launch, 100 by 25 Oct
Candidate signups	Track daily; watch the Brevo cap
Applications per real job	At least 5
Employer response time to applicants	Under 72 hours
Job list p95 latency / error rate	Under 1.5 s / under 1%

10. How we'll make money

Job portals do make money: Naukri earns almost all of its revenue from **employers**. Our edge is that we've already built an ATS pipeline, interview scheduling, scorecards, offers, assessments and AI interviews, while most small Indian companies still hire through Excel and WhatsApp.

Stream	Who pays	Notes
1. Freemium employer plan	Employers	Free: 1–2 active jobs and a basic pipeline. Paid: unlimited jobs, full ATS, interview and offer tools, team seats. The best long-term option.
2. Featured or boosted jobs	Employers	Pin a job to the top, or feature it in the WhatsApp channel and email. The easiest first sale; invoice by hand over UPI at first.
3. Candidate database access	Employers	Candidate search already exists. Charge per contact or per unlocked profile. Candidates must opt in.
4. AI screening	Employers	Assessments and AI interviews come back as a pay-per-candidate feature.
5. University placement portal	Colleges	A placement dashboard paid per student per year. Geeta University is the pilot.
6. Candidate premium	Candidates	AI tailored resumes (about ₹2.3 AI cost each), extra ATS checks, AI mock interviews, sold as small credit packs. Never charge candidates to apply.
7. Affiliate income	Partners	Pay-per-click on outbound jobs (check Careerjet, Jooble, Adzuna). Link skill gaps to Coursera or Udemy affiliate courses instead of running our own LMS.

Revenue roadmap

When	Step
11–25 Oct	Everything free. Build supply and trust. Ram interviews 10 employers about pricing (L04).
From about week 3	Sell featured jobs by hand (UPI invoice)
Month 2	Paid university placement pilot
Month 3	Paid employer plans; bring back assessments and AI interviews as paid features

Before we take money: register a business entity (a sole proprietorship is enough for Razorpay), fix the payment code properly, and move the frontend off Vercel Hobby.

Appendix A: Credential rotation runbook (Imran, 26–27 Sep)

Why: `server/.env` was committed on 23 Aug (commits `a110544` and `3214bdc`) and is still in git history, with the MongoDB connection string and `JWT_SECRET`. Razorpay test keys are also hardcoded in `paymentController.js`. Treat all of them as leaked.

Order matters: create the new value, update Render, redeploy, confirm the app works, and only then revoke the old value.

1. **MongoDB Atlas** (whoever holds the Atlas login): create a new database user or change the password. Set `MONGODB_URI` on Render and redeploy. Check that `/health` returns 200 and login works, then delete the old user. Note whether Network Access allows `0.0.0.0/0`.
2. **JWT_SECRET, SESSION_SECRET and OTP_HASH_SECRET**: generate a separate value for each with `node -e "console.log(require('crypto').randomBytes(48).toString('hex'))"` and set them on Render. Rotating the JWT secret logs everyone out, which is fine before launch.
3. **Razorpay**: regenerate the test-mode key in the dashboard. Task S02 removes the hardcoded fallbacks.
4. **Firebase web key and reCAPTCHA** (public by design): restrict the key to our domains under HTTP referrers, and keep only our domains on the reCAPTCHA site.
5. **Also check** `BREVO_API_KEY`, `CLOUDINARY_API_SECRET`, `FIREBASE_PRIVATE_KEY` and `GEMINI_API_KEY`. If any was ever committed, rotate it.
6. **Optional:** clean git history with `git filter-repo` after rotating. This rewrites history for everyone, so coordinate first. Rotation is what actually protects us.

Log the date and the person who did it in `docs/security-hardening.md`.

Appendix B: After launch (not in these 15 days)

- One auth model: HttpOnly cookie only, and remove the `localStorage` token.
- Move scheduled jobs out of the API into a separate worker, so the API can run as 2+ copies.
- Merge the three candidate profiles (student, fresher, professional) into one shared, configurable set of components.
- Split the largest controllers into service modules, and give dashboard tabs real URLs.
- Frontend tests, a lint gate in CI, structured logs and request IDs.
- Merge the parked branches (assessments, AI interviews, messaging) after their known issues are fixed.