

Sneha: your launch tasks

Candidate frontend owner. Eight tasks between 27 Sep and 5 Oct, with the details you need to start straight away.

8 tasks

4 of them P0 (must ship)

5 Oct

Last due date

6 Oct

Code freeze (evening)

11 Oct

Public launch

From	Ram Mohan Dixit, Manager
Date	26 September 2026
Branch	Start every task from <code>origin/launch/job-portal</code> , never from the <code>sneha</code> branch, which is 58 commits behind <code>main</code> .
Full team plan	<code>output/pdf/CareerConnect_Job_Portal_Launch_Plan.pdf</code> and <code>docs/launch/Launch_Tracker.csv</code>

Your role in one line: make the job portal **public, shareable, trustworthy and crash-proof** on the candidate side, and cleanly switch off everything that isn't launching.

1. Why these tasks, and what happens to your course work

For this launch the team decided to ship **only the job portal**, on **Sunday 11 October**. Most of your recent work (courses, LMS, course content upload, the student course pages, payment changes) is in the part we're **not** launching yet.

Your course work is not wasted. It stays in the code, switched off by a flag (task S01), and comes back after launch. The course code on your `sneha` branch is parked, not deleted. Courses tie into our revenue plan later (skill-gap courses, paid learning, and possibly university programmes).

For the next 9 days you own the **candidate-facing frontend of the job portal**: the pages visitors and candidates see first. These tasks decide whether people can find a job, share it and trust us, so they matter a lot.

Your tasks at a glance

ID	Pri	Task	Due	Reviewer	Depends on
B03	P1	Sneha branch (course page + course controller changes): out of launch scope; cherry-pick only the authMiddleware fix that treats literal 'null'/'undefined' Bearer tokens as missing	Mon 28 Sep	Vivek	—
G12	P1	Fix public-path 401 redirect list (jobs/internships/organizations) in client api.jsx	Tue 29 Sep	Tripti	—
S01	P0	Feature flag: hide courses, payments, assessments, AI interview (client VITE_ENABLE_* + server conditional route mounts)	Tue 29 Sep	Vivek	—
G06	P0	Global React error boundary + lazy-chunk load retry	Thu 01 Oct	Tripti	—
G01	P0	Public job and internship detail pages at /jobs/:id and /internships/:id (no login to view), fix employer /jobs/:id/applications 404	Sat 03 Oct	Tripti	S03
G05	P0	Privacy Policy, Terms, Contact pages; fix footer links; consent checkbox on all signup paths (DPDP notice + grievance contact)	Sun 04 Oct	Ram	L01
G02	P1	SEO: per-page title/description/OG tags, JobPosting JSON-LD on employer-posted jobs, sitemap.xml, robots.txt	Mon 05 Oct	Ram	G01
G11	P1	Report this job button + endpoint feeding admin reports	Mon 05 Oct	Vivek	—

Suggested day-by-day schedule

Day	Focus
Sun 27 Sep	Kickoff call at 10:00. Set up the branch. Finish B03 (small). Start S01 on the client (feature flag file, App.jsx routes).
Mon 28 Sep	S01 : sidebars, dashboards, Home.jsx. Open the B03 PR.
Tue 29 Sep	S01 : server mounts with Vivek, test all roles, open the PR. G12 (small).
Wed 30 Sep	Fix S01 review comments. Start G06 . Send Ram the page layout for the legal pages.
Thu 1 Oct	Finish G06 . Start G01 (JobDetailPage). Create the G05 page shells.
Fri 2 Oct	G01 : apply and redirect, sharing, public internship page. Ram's legal text is due today.
Sat 3 Oct	Finish G01 (employer 404 fix, mobile pass). G05 consent checkbox on the 4 signup paths.
Sun 4 Oct	Finish G05 (server consent with Vivek). Start G02 .
Mon 5 Oct	Finish G02 and G11 . Every PR open and reviewed.
6–10 Oct	Code freeze on the evening of 6 Oct, then only fix beta bugs Tripti assigns you. Do a mobile-layout pass on the pages candidates use most.

Other responsibilities

- **Daily employer outreach (J01):** 1 hour and 10 personal messages a day, like everyone on the team. Share your results in the evening update.
- **Review Ram's L02** (launch channels) from a candidate's point of view.
- **After launch (Q07):** triage at 10:00 each day and help fix candidate-side bugs until 25 Oct.

Who to talk to

Topic	Person
Any server change (app.js mounts, User consent field, report endpoint, sitemap, which routes are public)	Vivek , who also deploys. Only Vivek merges into <code>main</code> .
Legal text, whether certifications stay, product questions	Ram
Stable IDs for external jobs (I04), Sentry (I08)	Imran
Test checklist and bug reports	Tripti
Account deletion (linked from the privacy page)	Yug

Rules for every task

- Branch from `origin/launch/job-portal`, named like `sneha/S01-feature-flags`. Open the PR into `launch/job-portal`, never into `main`.
- PR title starts with the task ID, for example `[G01] Public job detail pages`. The description says what changed and how you tested it, with screenshots for UI.
- Before opening a PR: `cd client && npm run build` passes, you see no new console errors, and you've checked the page at 360 px mobile width.
- Never commit `.env` or `.env.local` files.
- Update your rows in `docs/launch/Launch_Tracker.csv` every evening: Not started, In progress, In review or Done.
- Blocked for more than 2 hours? Say so in the team group. Don't wait for the next standup.

B03 Bring your auth fix onto the launch branch

Priority	P1
Dates	Start Sun 27 Sep → due Mon 28 Sep
Reviewer	Vivek (code review), then Tripti tests it
Depends on	Nothing, so you can start right away
Done when	Auth fix on launch/job-portal with a test; course changes parked

Why

Your branch `sneha` has a useful fix in `server/middleware/authMiddleware.js` (commit `57a4711`). When the frontend has no token it sometimes sends the literal text `"null"` or `"undefined"` as a Bearer token, and the server tried to verify it as a real JWT. The rest of that branch is course work, which isn't launching, so we take only this fix.

Steps

1. Create a branch from the launch branch:

```
git fetch origin
git switch -c sneha/B03-auth-null-token origin/launch/job-portal
```

2. In `protect()` in `server/middleware/authMiddleware.js`, re-apply your 3 lines right after the token is read:

```
if (token === "null" || token === "undefined") {
  token = null;
}
```

Re-type them rather than cherry-picking the whole commit, because that commit also changes course files.

3. Add a unit test in `server/__tests__/unit/middleware/`, next to the existing auth middleware tests: a request with `Authorization: Bearer null` should get the normal "not authorized" 401, not a JWT error or a 500.
4. Run `cd server && npm run test:unit`, then open a PR titled `[B03] Treat "null"/"undefined" bearer tokens as missing` into `launch/job-portal`.

S01 Switch off courses, payments, assessments and AI interviews

Priority	P0
Dates	Start Sun 27 Sep → due Tue 29 Sep
Reviewer	Vivek (code review), then Tripti tests it
Depends on	Nothing, so you can start right away
Done when	No course/payment/assessment UI anywhere for any role; /api/courses, /api/course-content, /api/payment, /api/employer/learning, /api/assessments return 404

Why

We launch only the job portal. Courses, payments, assessments and AI interviews must be **invisible for every role**: no menu item, no home page banner, no working URL, no working API. **Nothing gets deleted**. Your course and LMS work stays in the code behind a flag and comes back after launch (see the last page). You know this code best, which is why this task is yours.

Step 1: one feature-flag file (client)

Create `client/src/config/features.js`. Everything defaults to **off**, so production is safe even if nobody sets an env var:

```
const on = (v) => String(v).toLowerCase() === "true";
export const FEATURES = {
  courses: on(import.meta.env.VITE_ENABLE_COURSES),
  payments: on(import.meta.env.VITE_ENABLE_PAYMENTS),
  assessments: on(import.meta.env.VITE_ENABLE_ASSESSMENTS),
  aiInterview: on(import.meta.env.VITE_ENABLE_AI_INTERVIEW),
};
```

For local testing, put `VITE_ENABLE_COURSES=true` in your own `client/.env.local`, and never commit it.

Step 2: every place to hide (checked on latest main)

File	What to hide
<code>client/src/App.jsx</code>	Routes <code>/courses</code> , <code>/courses/:id</code> , <code>/courses/:id/learn</code> (around lines 365–370), <code>/courses</code> , <code>/courses/:id</code> , <code>/my-courses</code> (377–379), and <code>/employer/courses*</code> (444–447). Wrap them in <code>{FEATURES.courses && (<>...</>)}</code> , so the URLs fall through to the 404 page.
<code>components/student-dashboard/Sidebar.jsx</code>	<code>courses</code> item (line 15). Also line 11, <code>certifications</code> : ask Ram whether it depends on courses.
<code>components/fresher-dashboard/FresherSidebar.jsx</code>	<code>courses</code> item (around line 32)
<code>components/professional-dashboard/ProfessionalSidebar.jsx</code>	<code>courses</code> (line 10), plus <code>certifications</code> (line 17) if Ram confirms
<code>components/employer/EmployerSidebar.jsx</code>	<code>assessments</code> "Recruitment Tests" (line 30) behind <code>FEATURES.assessments</code> . The whole "LEARNING & ORGANIZATION" group (lines 36–42: <code>courses</code> , <code>training</code> , <code>skill gaps</code>) behind <code>FEATURES.courses</code> . Rename "Hiring & Training Analytics" (line 48) to "Hiring Analytics".

<code>pages/student/StudentDashboard.jsx</code>	CourseRecommendationsCard (import line 26), the courses tab handling (124), the recommended-courses filter (245–248), the courses tab render (429)
<code>pages/fresher/FresherDashboard.jsx</code>	FresherSkillDevelopment and FresherRecommendedCourses (imports 15–16; renders around 231, 242, 363), courses tab (98, 323)
<code>pages/professional/ProfessionalDashboard.jsx</code>	Courses tab (152, 505)
<code>pages/employer/EmployerDashboard.jsx</code>	Courses, manage-courses and learning tabs (around line 1324). Also skip their data fetches when the flag is off, so the dashboard doesn't call disabled APIs.
<code>pages/Home.jsx</code>	Nav links to /courses (lines 136 and 323), the courses banner (650), the "Certification courses" section (848–885)
Also search for	<code>grep -rn "courses\ razorpay\ Assessment" client/src</code> . Look for dashboard widgets and quick-action cards that link to courses, such as QuickActionsCard.

Step 3: switch off the server routes (Vivek reviews)

In `server/app.js`, wrap these mounts so they're registered only when the matching flag is true:

- Lines 121–123: `/api/courses` (x2) and `/api/course-content`
- Line 124: `/api/payment`
- Line 134: `/api/employer/learning`
- Line 136: `/api/assessments`

```
const flag = (name) => String(process.env[name]).toLowerCase() === "true";
if (flag("ENABLE_COURSES")) {
  app.use("/api/courses", courseRoutes);
  app.use("/api/courses", courseContentRoutes);
  app.use("/api/course-content", courseContentRoutes);
  app.use("/api/employer/learning", employerLearningRoutes);
}
if (flag("ENABLE_PAYMENTS")) app.use("/api/payment", paymentRoutes);
if (flag("ENABLE_ASSESSMENTS")) app.use("/api/assessments", assessmentRoutes);
```

The student dashboard API also returns `recommendedCourses` (`server/controllers/studentController.js`). Return an empty list when courses are off. Tell Vivek before you change his files, and keep that change small.

How to test

1. With no `.env.local`, log in as a student, fresher, professional and employer. Check that no course, training, assessment or payment item appears anywhere, including the home page when logged out.
2. Typing `/courses`, `/my-courses` or `/employer/courses` in the address bar shows the 404 page.
3. `curl -i http://localhost:5001/api/courses` and `/api/payment/...` return 404.
4. With the flags set to true locally, everything comes back and still works.
5. The browser console shows no errors from dashboards calling APIs that are switched off.

G12 Stop bouncing visitors to /login on public pages

Priority	P1
Dates	Start Tue 29 Sep → due Tue 29 Sep
Reviewer	Tripti (code review), then Tripti tests it
Depends on	Nothing, so you can start right away
Done when	Logged-out visitor never bounced to /login on public pages

Why

In `client/src/api/api.jsx` (lines 47–54), a 401 response sends the user to `/login` unless the current page is on a short public list. `/jobs`, `/internships`, `/organizations` and the new pages aren't on it, so a logged-out visitor browsing jobs can suddenly land on the login page.

Steps

1. Replace the chain of `startsWith` calls with one list:

```
const PUBLIC_PREFIXES = ["/home", "/login", "/admin/login", "/register", "/forgot-password",  
  "/set-password", "/companies", "/opportunities", "/jobs", "/internships", "/organizations",  
  "/privacy", "/terms", "/contact", "/about", "/resume-builder"];  
const isPublicPage = currentPath === "/" ||  
  PUBLIC_PREFIXES.some((p) => currentPath === p || currentPath.startsWith(p + "/"));
```

Don't add `"/"` as a prefix, because every path starts with `/`.

2. Check the route list in `App.jsx` for any other page that logged-out users can open.

How to test

In a private window (logged out), open `/jobs`, `/jobs/in/delhi`, `/internships` and a company page. Click around for 2 minutes and confirm you're never sent to `/login`, except when you click Apply, which should still ask you to log in.

G06 Error boundary and retrying failed page loads

Priority	P0
Dates	Start Thu 01 Oct → due Thu 01 Oct
Reviewer	Tripti (code review), then Tripti tests it
Depends on	Nothing, so you can start right away
Done when	Thrown error shows fallback UI instead of blank page

Why

There's no React error boundary anywhere. If one component throws, the whole app goes blank. Also, all 52 pages are loaded on demand (`lazy()` in `App.jsx`). After every Vercel deploy the old file names disappear, so users with an open tab get "Failed to fetch dynamically imported module" and a blank screen. This will happen to real users every time Vivek deploys.

Steps

1. Create `client/src/components/common/ErrorBoundary.jsx` as a class component with `getDerivedStateFromError` and `componentDidCatch` . The fallback screen says "Something went wrong" and has **Reload** and **Go to home** buttons, styled like `NotFound.jsx` . Log the error with `console.error` for now; Imran's Sentry task (I08) plugs in here later.
2. In `App.jsx` , wrap `<Suspense>...</Suspense>` in `<ErrorBoundary>` . Also reset the boundary when the route changes (use `location.pathname` as its `key`), so moving to another page recovers.
3. Add a `lazyWithRetry` helper and use it instead of `lazy` for all pages:

```
export const lazyWithRetry = (importer) => lazy(async () => {
  try {
    const mod = await importer();
    sessionStorage.removeItem("cc_chunk_reload");
    return mod;
  } catch (err) {
    if (!sessionStorage.getItem("cc_chunk_reload")) {
      sessionStorage.setItem("cc_chunk_reload", "1");
      window.location.reload(); // pick up the new deploy once
      return new Promise(() => {});
    }
    throw err; // second failure: ErrorBoundary shows fallback
  }
});
```

How to test

- Temporarily add `throw new Error("test")` in one page, and check that you see the fallback, not a blank page.
- Run `npm run build && npm run preview`, open the site, rebuild in another terminal (the file names change), then navigate to a page you haven't opened yet. It should reload once and work.

G01 Public job and internship detail pages

Priority	P0
Dates	Start Thu 01 Oct → due Sat 03 Oct
Reviewer	Tripti (code review), then Tripti tests it
Depends on	S03 (S03 means only approved jobs are public). External jobs need Imran's I04 for stable IDs, due 4 Oct.
Done when	Shareable link opens for logged-out visitor; apply prompts login

Why

This is the most important task you own. Right now **a job has no page with its own link**. Job cards only offer Quick Apply or an external link (`pages/jobs/JobDiscoveryPage.jsx`). Without a link, nobody can share a job on WhatsApp or LinkedIn, Google can't index it, and our growth plan doesn't work. The internship detail page exists, but only for logged-in candidates.

The backend is ready

- GET `/api/jobs/:id` is already public (`server/routes/jobRoutes.js:15`).
- GET `/api/internships/:id` works without login (`optionalAuth` , `server/routes/internshipRoutes.js:83`).
- Ask Vivek to confirm that both return **404 for jobs that aren't Published** once S03 is merged. Pending jobs must never be public.

Steps

1. **New page** `client/src/pages/jobs/JobDetailPage.jsx` at route `/jobs/:id` in the **public** part of `App.jsx` , next to the other `/jobs/*` routes (around line 298). React Router ranks fixed paths like `/jobs/latest` above `:id` , so they keep working. Show: title, company (linking to the public company page), location and work mode, salary **only if the employer entered one**, experience, skills, description, responsibilities, deadline and posted date.
2. **Apply button**: if logged in as a candidate, use the existing apply flow. If logged out, go to `/login?redirect=/jobs/<id>` (the same pattern as `JobDiscoveryPage.jsx:110`) and bring the user back afterwards. Employers and admins don't see Apply.
3. **Share button**: "Copy link" plus WhatsApp (`https://wa.me/?text=` + an encoded title and URL).
4. **Other states**: a loading skeleton, "This job is closed or no longer available" for 404 or expired jobs (with a link back to `/jobs`), and an error message with a retry button.
5. **Job cards**: in `JobDiscoveryPage.jsx` and the dashboard job cards, make the title and card open `/jobs/:id` . Keep Quick Apply as it is.
6. **Internships**: `/internships/:id` is currently inside the candidate-only `RoleProtectedRoute` block (`App.jsx` around lines 363 and 372; it's listed twice). Move it to the public routes, and make `pages/student/InternshipDetail.jsx` work when logged out, with the same Apply and redirect behaviour.
7. **External jobs** (`isExternal` or `scraped-*` IDs): until Imran's I04 is done, their IDs change on every request. Link these cards straight to the source site, and don't show a detail page for them yet. After I04 lands, show a detail page with an "Apply on <source>" button and a "Source: ..." credit line.
8. **Fix the employer 404**: `client/src/features/employer/dashboard/DashboardPage.jsx:179` navigates to `/jobs/${jobId}/applications` , which doesn't exist. Change it to open the employer dashboard's applicants (ATS) view filtered to that job; ask Vivek which tab and prop to use.
9. **Mobile first**: most users will open these links from WhatsApp on a phone. Test at 360 px width.

How to test

- Logged out, paste a job link into a new private window. The page loads and shows full details.
- Click Apply, log in, and check that you come back to the same job and can apply.
- A pending job's link shows "not available". A random ID shows "not available", not a crash.
- The employer dashboard "view applications" button no longer shows 404.

G05 Privacy Policy, Terms, Contact, and consent at signup

Priority	P0
Dates	Start Thu 01 Oct → due Sun 04 Oct
Reviewer	Ram (code review), then Tripti tests it
Depends on	L01. Ram must give you the final wording by 2 Oct.
Done when	Pages live; signup blocked without consent; consent timestamp stored

Why

We collect resumes and contact details, which are personal data under India's **DPDP Act 2023**. Right now the footer links "Privacy Policy", "Terms & Conditions", "About" and "Sitemap" all point to `/home` (`pages/Home.jsx` lines 1016, 1028, 1030, 1032), and **no signup page asks for consent**. We can't launch like this.

Steps

1. **Pages:** create `client/src/pages/legal/PrivacyPolicy.jsx`, `Terms.jsx` and `Contact.jsx`, with public routes `/privacy`, `/terms` and `/contact`. Build a simple readable layout now with placeholder headings, and drop in Ram's text when it arrives. Show a "Last updated" date and a version string such as `2026-10-01`.
2. **What the privacy page must cover** (Ram writes it, you check that nothing is missing): what data we collect (profile, resume, contact details), why (showing your profile to employers *you apply to*, and job matching), who we share it with, how long we keep it, your rights (access, correction, deletion through account deletion, which is Yug's G07), cookies, and a **grievance contact** (name and email).
3. **What the terms page must cover: employers must never charge candidates any fee**, what employers must provide, content we don't allow, and that we can remove jobs and accounts.
4. **Footer:** point the `Home.jsx` links at `/privacy`, `/terms` and `/contact` (use `/contact` for About for now), and point Sitemap at `/sitemap.xml` after G02.
5. **Consent checkbox** on **all four** signup paths: `pages/auth/Signup.jsx`, `EmployerRegister.jsx`, `GoogleOnboarding.jsx` and `GoogleEmployerOnboarding.jsx`. The label reads: "I agree to the Terms and Privacy Policy", with links that open in a new tab. The submit button stays disabled until it's ticked. Send `acceptedTerms: true` and `termsVersion` with the request.
6. **Server (Vivek reviews):** add `consent: { termsVersion: String, acceptedAt: Date }` to `server/models/User.js`, and make the register and onboarding endpoints in `authController.js` **reject** requests without `acceptedTerms`. Checking only in the UI isn't enough.

How to test

- All four signup paths can't be submitted without ticking the box. A direct API call without `acceptedTerms` returns 400.
- A new user in MongoDB has `consent.acceptedAt` set.
- The footer links open the right pages when logged out.

G02 SEO: page titles, previews, JobPosting data, sitemap

Priority	P1
Dates	Start Fri 02 Oct → due Mon 05 Oct
Reviewer	Ram (code review), then Tripti tests it
Depends on	G01
Done when	Google Rich Results Test passes for a job page; sitemap lists live jobs

Why

This is free traffic. Google for Jobs shows jobs directly in search results when a page has **JobPosting structured data**. WhatsApp and LinkedIn show a preview card only if the page has Open Graph (OG) tags. Today `client/index.html` has just one title and a favicon, and there's no `robots.txt` or `sitemap.xml`.

Steps

1. **Default tags in `client/index.html`**: meta description, `og:title`, `og:description`, `og:image` and `twitter:card`. Make a 1200×630 OG image and save it as `client/public/og-image.png`.
2. **Per-page tags**: we're on React 19, which lets a component render `<title>` and `<meta>` directly; React moves them into `<head>`. No library needed. On `JobDetailPage`, use "`<Job title> at <Company>, <City> | CareerConnect`" with a 160-character description. Do the same for internship, company and job list pages.
3. **JobPosting JSON-LD on `JobDetailPage`**, **only for jobs posted by employers on our site**. Never add it to external or scraped jobs, because Google penalises that.

```
<script type="application/ld+json">{JSON.stringify({
  "@context": "https://schema.org/", "@type": "JobPosting",
  title, description /* HTML allowed */, datePosted: createdAt,
  validThrough: deadline, employmentType /* FULL_TIME, INTERN... */,
  hiringOrganization: { "@type": "Organization", name: companyName, sameAs: companyWebsite },
  jobLocation: { "@type": "Place", address: { "@type": "PostalAddress",
    addressLocality: city, addressRegion: state, addressCountry: "IN" } },
  // remote jobs: jobLocationType: "TELECOMMUTE" + applicantLocationRequirements
  // baseSalary ONLY if the employer entered a real salary
}}}</script>
```

4. **`client/public/robots.txt`**: Allow: `/`; Disallow: `for /admin, /employer, dashboards and /applications;` plus a Sitemap: `https://<our-domain>/sitemap.xml` line.
5. **`sitemap.xml`** must list live jobs, so it has to come from the server. Ask Vivek or Imran for a small `GET /sitemap.xml` endpoint that lists published job, internship and company URLs. Then add a rewrite in `client/vercel.json` **above** the catch-all `/(.*) → /index.html` rule, pointing `/sitemap.xml` at the Render API. Vivek deploys it.

How to test

- Paste a live job URL into Google's **Rich Results Test**. It should show "Job posting: valid".
- Share a job link in a WhatsApp chat with yourself. The preview card should show a title and image.
- `/robots.txt` and `/sitemap.xml` load on the production domain.
- Our site is a single-page app, so Google needs extra time to render it. Ram submits the sitemap in Google Search Console after launch.

G11 "Report this job" button

Priority	P1
Dates	Start Sat 03 Oct → due Mon 05 Oct
Reviewer	Vivek (code review), then Tripti tests it
Depends on	Nothing, so you can start right away
Done when	Report appears in admin reports within one minute

Why

Fake jobs that ask candidates for a "registration fee" are the biggest trust risk for Indian job portals. Users must be able to flag a job, and admins must see it quickly. The `Report` model already supports this (`server/models/Report.js`: categories "Opportunity" and "Spam / Fraud", with status and priority). But the only create endpoint is `POST /api/admin/reports` (`adminRoutes.js:164`), which is for admins only.

Steps

1. **Server (agree the endpoint with Vivek first):** a user endpoint such as `POST /api/reports` behind `protect`, with a rate limit of about 5 per hour per user. The body is `{ opportunityType, opportunityId, reason, details }`. It saves a `Report` with category "Spam / Fraud" or "Opportunity", **priority High when the reason is "asks for money"**, and the reporter's ID.
2. **UI:** a small "Report" link on `JobDetailPage` and the internship detail page opens a modal. The reason is a dropdown: "Asks candidates for money", "Fake or scam job", "Wrong or misleading information", "Duplicate", "Other". There's an optional details box of up to 500 characters. Logged-out users are asked to log in first. After submitting, show "Thanks, our team will review this within 24 hours."
3. Admins already have `/admin/reports`. Check that the new reports appear there.

How to test

Report a job as a candidate, then log in as admin. The report shows on `/admin/reports` within a minute, and "asks for money" reports show as High priority. A sixth report within an hour gets a friendly 429 message.

Your checklist before the go/no-go meeting (10 Oct)

- ☐ No course, payment, assessment or AI interview UI is visible for any role, and those URLs and APIs return 404 (S01).
- ☐ A logged-out visitor can open a shared job or internship link, see full details, and is brought back after logging in to apply (G01).
- ☐ The employer "view applications" link no longer shows 404 (G01).
- ☐ Logged-out visitors are never sent to /login while browsing public pages (G12).
- ☐ A crash in any page shows the fallback screen, and a new deploy doesn't blank open tabs (G06).
- ☐ Privacy, Terms and Contact pages are live; all 4 signup paths require consent, and the server stores it (G05).
- ☐ A job URL passes Google's Rich Results Test; WhatsApp shows a preview; robots.txt and sitemap.xml work (G02).
- ☐ Users can report a job, and it shows up in /admin/reports (G11).
- ☐ Every one of your PRs is merged into `launch/job-portal`, and your tracker rows say Done.

After launch: once the job portal is stable, your course and LMS work comes back as a planned feature, with a proper payment flow (the current payment code needs security fixes first, which is part of why it's off for launch).